

Hesitate — Demo Video Script

YC Fall 2026 × Moss — Zero Latency Builder Sprint. Total runtime: ~2 minutes 40 seconds.

Before you press record

Open Terminal. Go to the project folder by typing this and pressing Enter:

```
cd "/Users/ashraf/Desktop/PROJECTS/YC Fall 2026 x Moss"
```

Turn on the project's Python environment and load your saved keys, by typing each of these and pressing Enter:

```
source .venv/bin/activate
```

```
set -a; source .env; set +a
```

Start the voice agent so it's ready to answer a call later in the video. Type this and press Enter, then leave this window open and running for the rest of the recording:

```
python -m agent.voice.entrypoint start
```

Open a second Terminal window (or tab) for everything else — don't close the one running the agent.

Now start your screen recorder (QuickTime: File → New Screen Recording, or any recorder you're used to), then begin the script below.

[0:00–0:12] Open cold, no title card

DO THIS	In your second Terminal window, just leave a plain empty prompt on screen — don't type anything yet. This is the very first thing viewers see.
CAPTION	"A clinic voice agent is about to make a mistake."
SAY THIS	<i>Say you call a clinic and ask how long you need to fast before a blood test. The clinic's actual policy says eight hours. But there's an old sheet still sitting in their system that says twelve. Watch what happens.</i>

[0:12–0:35] Run the real scenario

DO THIS	Type this one command and press Enter. Real output will print on screen — let it run, don't skip past it: <code>python -m pytest tests/test_end_to_end_live.py -v -s</code>
CAPTION	"Real retrieval. Real model. Nothing staged."
SAY THIS	<i>This isn't a script I wrote to look bad. I asked a real search system — Moss — to pull up the clinic's fasting policy, and it actually ranks the old, wrong sheet slightly higher than the current one. Then I handed that to a real language model and asked it the same question a patient would ask.</i>

[0:35–0:55] The model gets it wrong

DO THIS	The command from the last step is still printing output. Find the line that starts with “LLM reply” — it will say something like “you'll need to fast for twelve hours.” Move your mouse cursor next to that line so viewers' eyes go there, or use your recorder's zoom/highlight tool if it has one.
CAPTION	“The model said twelve hours. That's wrong.”
SAY THIS	<i>And there it is — it tells you twelve hours, with total confidence. If this were a normal voice agent, that's what the patient hears. They skip breakfast for half a day longer than they need to, or worse, they show up having eaten and the appointment gets cancelled.</i>

[0:55–1:15] The catch

DO THIS	Same terminal, scroll down slightly if needed to the line that says CONTRADICTED, and the corrected sentence right after it (something like “Actually — 8 hours, per your prep instructions”). Point your cursor at it or read it out loud slowly so the viewer can follow along.
CAPTION	“Caught and corrected — before anyone hears it.”
SAY THIS	<i>Before that sentence ever reaches speech, my system checks it against the clinic's actual current policy. It sees the mismatch, throws out the wrong number, and replaces it with the real one — eight hours. That whole check happens in about ten milliseconds. Fast enough that it never breaks the flow of the conversation.</i>

[1:15–1:35] Why this is different from normal AI search

DO THIS	Switch away from the terminal. Open the architecture diagram image you got from the HiDevs Architecture Copilot (save it as a PNG or PDF beforehand, then just open it in Preview). Let it sit on screen for the whole 20 seconds — don't scroll around on it, just let people look at it while you talk.
CAPTION	“Most AI agents check the input. This checks the output.”
SAY THIS	<i>Most systems like this check what the model was given before it answers. That's not the same as checking what it actually says. A model can be handed the right information and still get the answer wrong. So instead of trusting it, I check every sentence it's about to speak, one at a time, against the real record — and if there's no clear answer, it just says so instead of guessing.</i>

[1:35–2:05] The live call

DO THIS	Open your web browser and go to this address: hesitate-v2.onrender.com/call.html — Click the button that says “Start call.” Wait a couple seconds for it to connect. Then speak clearly into your microphone: “How long do I need to fast before my bloodwork?” Stay quiet and let the agent’s voice answer — don’t talk over it.
CAPTION	“A real call, right now, no editing.”
SAY THIS	<i>(Before clicking Start call) “This is the actual product, running live, right now. I’m going to call it myself.”</i> <i>(Then stay quiet while the real exchange plays out.)</i>

[2:05–2:25] It's honestly real, not a mockup

DO THIS	Switch browser tabs three or four times, staying on each for about 3 seconds: 1) your GitHub repo page showing the commit history, 2) your Moss project dashboard, 3) the deployed link hesitate-v2.onrender.com loading in the browser.
CAPTION	“Every account behind this is real.”
SAY THIS	<i>Everything you just saw is running on real accounts — a real search backend, a real phone system, a real hosted server. The code is public, including the mistakes I found and fixed along the way. I didn’t hide those, because the whole point of this project is catching mistakes instead of hiding them.</i>

[2:25–2:40] Honest limits, then close

DO THIS	Turn your webcam on and just look at the camera and talk for this part, or keep any calm screen showing (like the homepage of your deployed link). Nothing needs to be clicked here.
CAPTION	“It doesn’t catch everything yet. We say so.”
SAY THIS	<i>It’s not finished. Right now it catches specific kinds of claims — numbers, dates, coverage, what to bring. It doesn’t catch every vague thing a model might say yet, and I’m upfront about that in the write-up. What made this possible at all is that the search step is fast enough to run in the middle of a live conversation instead of after it. Repo and live link are below — thanks for watching.</i>

If something goes wrong while recording

- Don’t rehearse it into stiffness — read it once for timing, then say it in your own words the second time. A little imperfection reads as more real, not less.

- If the live call in the 1:35–2:05 step doesn't get a clean spoken answer, don't fake it. Instead, open hesitate-v2.onrender.com/gate/live-demo in your browser and show that real output, and change the line to: “here's the same check running against the real Moss service.”
- If you make a mistake mid-sentence, just pause, take a breath, and say the sentence again — you can cut the pause out later, or just leave it in. Small imperfections are fine.